

Introduction to Deep Learning (I2DL)

Mock Exam

IN2346 - SoSe 2020

Technical University of Munich

Problem		Full Points	Your Score
1	Multiple Choice	14	
2	Short Questions	16	
3	Logistic Regression	11	
4	Backpropagation	9	
5	Layers	12	
6	CNN Architectures	9	
Total		71	

Total Time: **71 Minutes**

Allowed Ressources: **None**

The purpose of this mock exam is to give you an idea of the type of problems and the structure of the final exam. The mock exam is not graded. The final exam will most probably be composed of 90 graded points with a total time of 90 minutes.

Multiple Choice Questions:

- For all multiple choice questions any number of answers, i.e. either zero (!) or one or multiple answers can be correct.
- For each question, you'll receive 2 points if all boxes are answered correctly (i.e. correct answers are checked, wrong answers are not checked) and 0 otherwise.

How to Check a Box:

- Please **cross** the respective box: ☒ (interpreted as **checked**)
- If you change your mind, please **fill** the box: ☐ (interpreted as **not checked**)
- If you change your mind again, please **circle** the box: ☐ (interpreted as **checked**)

Part I: Multiple Choice (14 points)

1. (2 points) To avoid overfitting, you can...
 - ☐ increase the size of the network.
 - ☐ use data augmentation.
 - ☐ use Xavier initialization.
 - ☐ stop training earlier.

2. (2 points) What is true about Dropout?
 - ☐ The training process is faster and more stable to initialization when using Dropout.
 - ☐ You should not use leaky ReLu as non-linearity when using Dropout.
 - ☐ Dropout acts as regularization.
 - ☐ Dropout is applied differently during training and testing.

3. (2 points) What is true about Batch Normalization?
 - ☐ Batch Normalization uses two trainable parameters that allow the network to undo the normalization effect of this layer if needed.
 - ☐ Batch Normalization makes the gradients more stable so that we can train deeper networks.
 - ☐ At test time, Batch Normalization uses a mean and variance computed on training samples to normalize the data.
 - ☐ Batch Normalization has learnable parameters.

4. (2 points) Which of the following optimization methods use first order momentum?
 - ☐ Stochastic Gradient Descent
 - ☐ Adam
 - ☐ RMSProp
 - ☐ Gauss-Newton

5. (2 points) Making your network deeper by adding more parametrized layers will always...
 - ☐ slow down training and inference speed.
 - ☐ reduce the training loss.
 - ☐ improve the performance on unseen data.
 - ☐ (Optional: make your model sound cooler when bragging about it at parties.)

6. (2 points) Which of the following are LSTM gates?

- ☐ Input Gate
- ☐ Output Gate
- ☐ Remember Gate
- ☐ Forget Gate

7. (2 points) What can 1×1 convolutions be used for?

- ☐ To perform feature selection.
- ☐ To reduce the number of channels before a costly operation.
- ☐ To make a network more complex as we can add another non-linearity after the convolution.
- ☐ To replace fully-connected layers in order to make a network fully-convolutional.

Part II: Short Questions (16 points)

1. (2 points) You're training a neural network and notice that the validation error is significantly lower than the training error. Name two possible reasons for this to happen.

2. (2 points) You're working for a cool tech startup that receives thousands of job applications every day, so you train a neural network to automate the entire hiring process. Your model automatically classifies resumes of candidates, and rejects or sends job offers to all candidates accordingly. Which of the following measures is more important for your model? Explain.


$$\text{Recall} = \frac{\text{True Positives}}{\text{Total Positive Samples}}$$

$$\text{Precision} = \frac{\text{True Positives}}{\text{Total Predicted Positive Samples}}$$

3. (2 points) You're training a neural network for image classification with a very large dataset. Your friend who studies mathematics suggests: "If you would use Newton-Method for optimization, your neural network would converge much faster than with gradient descent!". Explain whether this statement is true (1p) and discuss potential downsides of following his suggestion (1p).

4. (2 points) Your colleague trained a neural network using standard stochastic gradient descent and L2 weight regularization with four different learning rates (shown below) and plotted the corresponding loss curves (also shown shown below). Unfortunately he forgot which curve belongs to which learning rate. Please assign each of the learning rate values below to the curve (A/B/C/D) it probably belongs to and explain your thoughts.

learning_rates = [$3e-4$, $4e-1$, $2e-5$, $8e-3$]



5. (2 points) Explain how and why LSTM networks often outperform traditional RNNs.

6. (1 point) Explain why we need activation functions.

7. (3 points) When implementing a neural network layer from scratch, we usually implement a 'forward' and a 'backward' function for each layer. Explain what these functions do, potential variables that they need to save, which arguments they take, and what they return.

8. (2 points) Optional: Given a Convolution Layer with 8 filters, a filter size of 6, a stride of 2, and a padding of 1. For an input feature map of $32 \times 32 \times 32$, what is the output dimensionality after applying the Convolution Layer to the input?

Part III: Logistic Regression (11 points)



1. (11 points) We want to build a binary classifier based on a logistic regression model that detects dogs, i.e. returns a probability p that a given sample is a dog, while the probability for not being a dog is $1 - p$.

To predict probabilities p , we use the Sigmoid Function $\sigma(z)$.

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- (a) (3 points) Show that the derivative of the Sigmoid Function $\sigma(z)$ w.r.t. the input z is equal to $\sigma(z)(1 - \sigma(z))$.

- (b) (3 points) Using the result above, compute the derivative of the Binary Cross Entropy Loss (of a single sample) w.r.t. the score z .




- (c) (5 points) Assume you are training our logistic regression classifier on a single picture of a dog ($t = 1$), with features $x = [0.8, 0.5, -1]^T$. The current model parameters are $w = [0.5, 0.4, -0.3]^T$ and $b = 0.2$. Perform a gradient descent iteration with learning rate $\alpha = 0.1$ and update the weight w_1 (indexing starts with 1) with respect to the Binary Cross Entropy Loss. You can use the plot above to estimate values of the sigmoid function.



Part IV: Backpropagation (9 points)

1. (9 points) Given the following neural network with fully connection layer and ReLU activations, including two input units (i_1, i_2), four hidden units (h_1, h_2) and (h_3, h_4). The output units are indicated as (o_1, o_2) and their targets are indicated as (t_1, t_2). The weights and bias of fully connected layer are called w and b with specific sub-descriptors.



The values of variables are given in the following table:

Variable	i_1	i_2	w_{11}	w_{12}	w_{21}	w_{22}	w_{31}	w_{32}	w_{41}	w_{42}	b_1	b_2	b_3	b_4	t_1	t_2
Value	2.0	-1.0	1.0	-0.5	0.5	-1.0	0.5	-1.0	-0.5	1.0	0.5	-0.5	-1.0	0.5	1.0	0.5

- (a) (3 points) Compute the output (o_1, o_2) with the input (i_1, i_2) and network parameters as specified above. Write down all calculations, including intermediate layer results.

- (b) (1 point) Compute the mean squared error of the output (o_1, o_2) calculated above and the target (t_1, t_2) .

- (c) (5 points) Update the weight w_{21} using gradient descent with learning rate 0.1 as well as the loss computed previously. (Please write down all your computations.)

Part V: Layers (12 points)

1. (8 points) You are given the following kernel W of a convolutional layer:

$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$

W


- (a) (2 points) Apply the kernel W with a stride of $s = 2$ and a zero padding of $p = 1$ to the input X below. Fill in the missing values (red boxes) of the output $Z = X * W$.

3	1	6	8	0
2	3	2	7	-6
-8	0	6	-5	-8
1	2	5	-2	-4
8	7	-2	-1	7

X

1		1
0		
	1	0

Z




- (b) (4 points) Given input X , and the gradients ∇Z of the output, calculate the gradients ∇W of the kernel. Fill in the missing values (red boxes) in the figure below.

1	1	1
1	0	1
1	1	1

 ∇Z

7		3
10		
8	-5	

 ∇W 

- (c) (2 points) Is this convolutional layer the same as a 3×3 average pooling layer? Do they compute the same function? Do they behave similarly during train and test?

2. (4 points) As you learned in the lecture, the use of skip connections in ResNets has substantially improved the Deep Learning state-of-the-art by allowing us to train much deeper networks.
- (a) (2 points) Draw a residual block together with a skip connection below. The residual block should contain two layers, which you can represent by l_1 and l_2 . For simplicity, you don't need to draw in ReLU-activations.




- (b) (1 point) For your above drawing, given the partial derivative of the residual block $R(x) = l_2(l_1(x))$ as $\frac{\partial R(x)}{\partial x} = r$, calculate $\frac{\partial H(x)}{\partial x}$.

- (c) (1 point) Why can we build deeper networks with residual blocks?

Part VI: CNN Architectures (9 points)

1. (9 points) After having obtained strong programming experience in I2DL, you want to demonstrate your acquired deep learning skills by competing in a renowned image classification contest. In this contest, you have to classify images of size 256×256 into one of 1000 classes. Since you learned that CNNs are great at such tasks, you decide to start out with the following CNN architecture:




Notes:

- The values directly below the arrows indicate the filter sizes f of the corresponding convolution and pooling operations.
- s stands for stride. If no stride is specified, a stride of $s = 1$ is used.
- All convolution and pooling layers use a padding of $p = \frac{f-1}{2}$, for corresponding filter size f .

- (a) (1 point) Which classic network architecture is similar to the architecture above?

- (b) (1 point) In the figure above, the output layer is missing. Explain which type of layer you would use there, what its dimensions would be, and which output and loss functions you would choose.

- (c) (3 points) Explain how to calculate output sizes of convolutional layers. Then, apply your formula to find the heights and widths of all layers in the network, and fill in your results in the figure below.



- (d) (1 point) Calculate the number of parameters in the first convolutional layer.

- (e) (1 point) If we apply this network architecture to larger input images, would the total amount of parameters in the network change? Justify your answer.

- (f) (2 points) Unfortunately, the dataset of the competition is rather small. Online you found a much larger image classification dataset containing almost similar images, but the set of classes is different. Explain how you could make use of the larger dataset, and how each layer in your network would be affected in the process.



Additional Space for solutions. Clearly mark the problem your answers are related to and strike out invalid solutions.

