

Introduction to Deep Learning (I2DL)

Exercise 7: Pytorch

Deep Learning Frameworks

The two big ones

- Tensorflow - Google
 - As well as Keras
- Pytorch - Facebook



Other examples

- CNTK - Microsoft
- Mxnet - Apache



Static Framework: Tensorflow

Network Declaration

```
N, D, H = 64, 1000, 100


x = tf.placeholder(tf.float32, shape=(N, D))
y = tf.placeholder(tf.float32, shape=(N, D))
w1 = tf.Variable(tf.random_normal((D, H)))
w2 = tf.Variable(tf.random_normal((H, D)))

h = tf.maximum(tf.matmul(x, w1), 0)
y_pred = tf.matmul(h, w2)
diff = y_pred - y
loss = tf.reduce_mean(tf.reduce_sum(diff ** 2, a
grad_w1, grad_w2 = tf.gradients(loss, [w1, w2])

learning_rate = 1e-5
new_w1 = w1.assign(w1 - learning_rate * grad_w1)
new_w2 = w2.assign(w2 - learning_rate * grad_w2)
updates = tf.group(new_w1, new_w2)
```

Iteration

```
with tf.Session() as sess:
    sess.run(tf.global_variables_initializer())
    values = {x: np.random.randn(N, D),
              y: np.random.randn(N, D),}
    losses = []
    for t in range(50):
        loss_val, _ = sess.run([loss, updates], feed_dict=values)
```



What if you want to change the learning rate?

Dynamic Framework: Pytorch

Network Declaration

```
from torch.autograd import Variable

N, D_in, H, D_out = 64, 1000, 100, 10

x = Variable(torch.randn(N, D_in).cuda(),
y = Variable(torch.randn(N, D_out).cuda(),
w1 = Variable(torch.randn(D_in, H).cuda(),
w2 = Variable(torch.randn(H, D_out).cuda())
```

Iteration

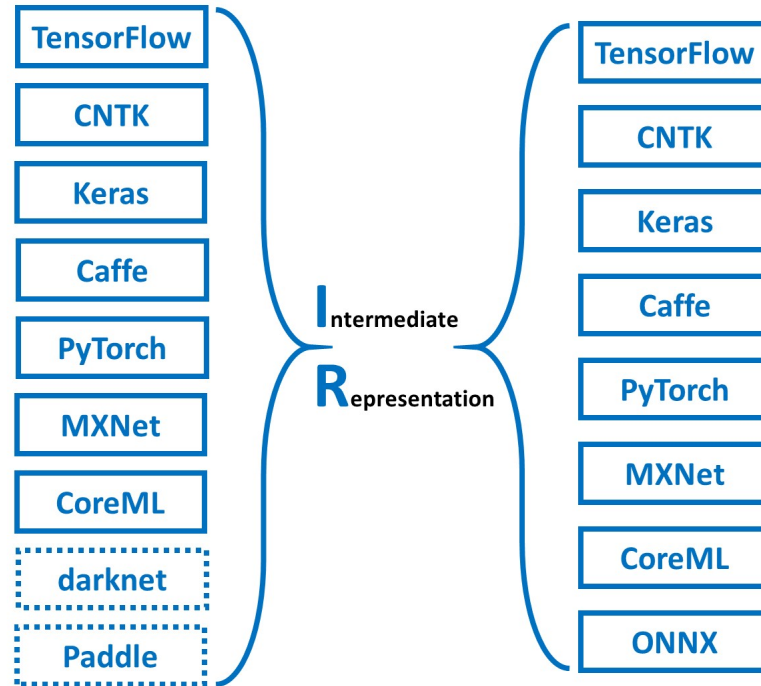
```
learning_rate = 1e-6
for t in range(500):
    y_pred = x.mm(w1).clamp(min=0).mm(w2)
    loss = (y_pred - y).pow(2).sum()

    if w1.grad: w1.grad.data.zero_()
    if w2.grad: w2.grad.data.zero_()
    loss.backward()

    w1.data -= learning_rate * w1.grad.data
    w2.data -= learning_rate * w2.grad.data
```

-> Advantages/Disadvantages

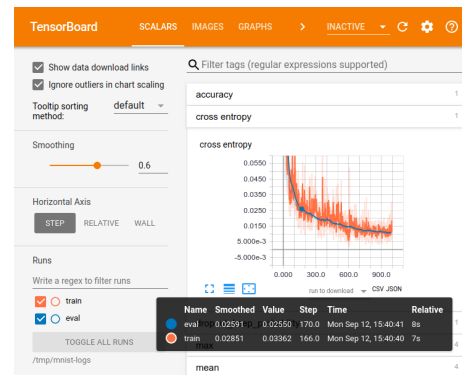
Framework Conversion



See: <https://github.com/microsoft/MMdnn>

Today's Contents

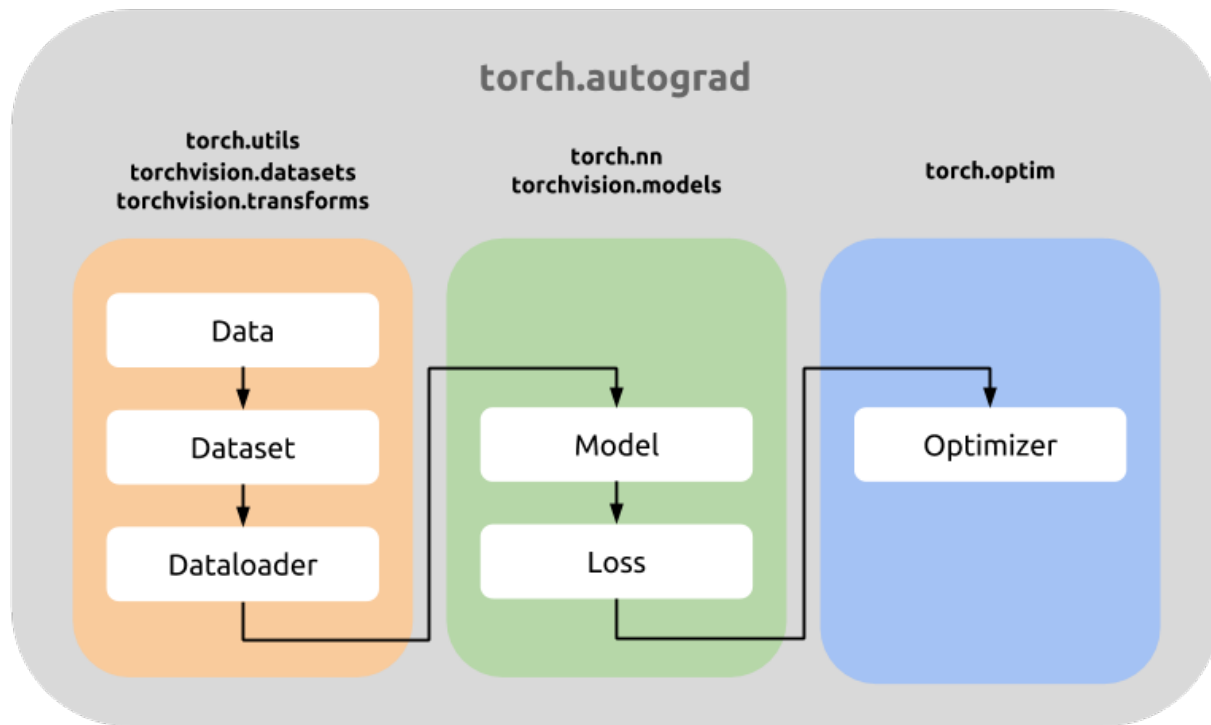
- Pytorch
 - Networks & Data
- Tensorboard
 - Easy visualization on the fly
 - Part of Tensorflow, but also accessible in Pytorch
- Pytorch Lightning
 - Solver



No submission
;-)

Pytorch

Pytorch: Overview



PyTorch vs. NumPy

- PyTorch shares lots of similarities with NumPy.
 - `np.array` vs. `torch.Tensor`

```
a_np = np.array([[1,2,3],[5,6,7]]) #NumPy array
a_ts = torch.tensor([[1,2,3],[4,5,6]]) # Tensor
print(a_np)
print(a_ts)
```

```
[[1 2 3]
 [5 6 7]]
tensor([[1, 2, 3],
        [4, 5, 6]])
```

PyTorch vs. NumPy

- Similar Operations

- Indexing

```
# indexing
b = a_ts[:2, :2] # use numpy type indexing
b[:, 0] = 0 # For assignment
b
```

```
tensor([[0, 2],
        [0, 5]])
```

- Mathematical operations
 - Etc.

Easy Device Assignment

```
device = torch.device("cuda:0" if torch.cuda.is_available() else "cpu")
print(device)

print(f"Original device: {x.device}") # "cpu", integer


tensor = x.to(device)
print(f"Current device: {x.device}") #"cpu" or "cuda", double
```

```
cpu
Original device: cpu
Current device: cpu
```

Datasets: Torchvision

- Torchvision
 - `torchvision.datasets` contains many datasets, such as ImageNet, FashionMNIST, etc.

```
fashion_mnist_dataset = torchvision.datasets.FashionMNIST(root='../datasets',  
                                                         train=True,  
                                                         download=True,  
                                                         transform=transform)  
  
fashion_mnist_dataloader = DataLoader(fashion_mnist_dataset, batch_size=8)
```



Source: https://www.tensorflow.org/datasets/catalog/fashion_mnist



Easy Network Creation

```
import torch.nn as nn
# defining the model
class Net(nn.Module):
    def __init__(self, input_size=1*28*28, output_size=100):
        super(Net, self).__init__()
        self.fc1 = nn.Linear(input_size, output_size)

    def forward(self, x):
        x = self.fc1(x)
        return x

net = Net()
net = net.to(device)
```

Where is the
backward
pass?



Training Loop

`zero_grad()`

`y_pred = net(X)`

`loss = criterion(y_pred, y)`

`loss.backward()`

`optimizer.step()`

zero the gradient

complete a forward pass

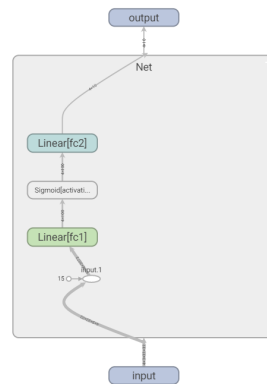
calculate the loss

backpropagation

apply gradients

```
# zero the parameter gradients
optimizer.zero_grad()

# forward + backward + optimize
y_pred = net(X) # input x and predict based on x
loss = criterion(y_pred, y) # calculate the loss
loss.backward() # backpropagation, compute gradients
optimizer.step() # apply gradients
```



References on Pytorch

- Documentation(1.4.0):
https://tutorials.pytorch.kr/beginner/nlp/pytorch_tutorial.html
- Repository: <https://github.com/pytorch/pytorch>
- Examples (very nice):
<https://github.com/pytorch/examples>
- PyTorch for NumPy users:
<https://github.com/wkentaro/pytorch-for-numpy-users>

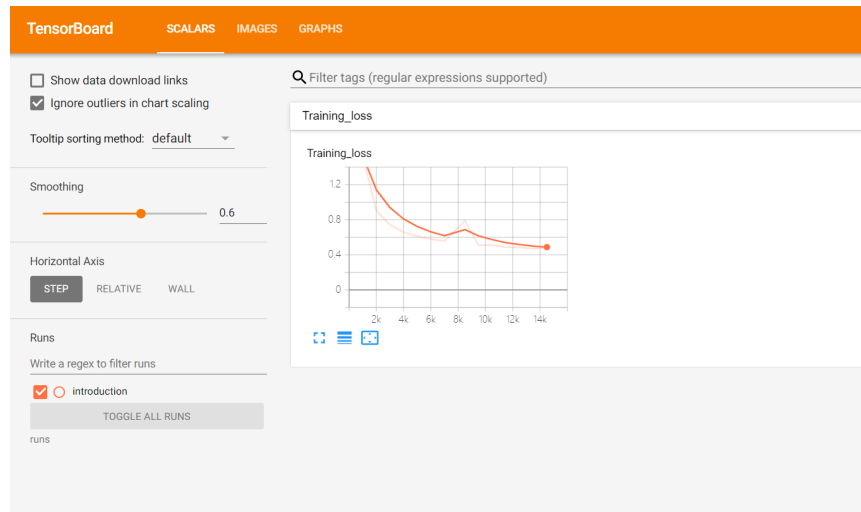
Tensorboard

Simple Logging

- Directly access tensorboard in your training loop

```
# ...log the running loss
writer.add_scalar('Training loss',
                 running_loss / 1000,
                 epoch * len(trainloader) + i)
```

- Tensorboard generates the graph/timestamps etc. for you



Visualize Network Architectures

- Using a single forward pass, tensorflow can map and display your network graph

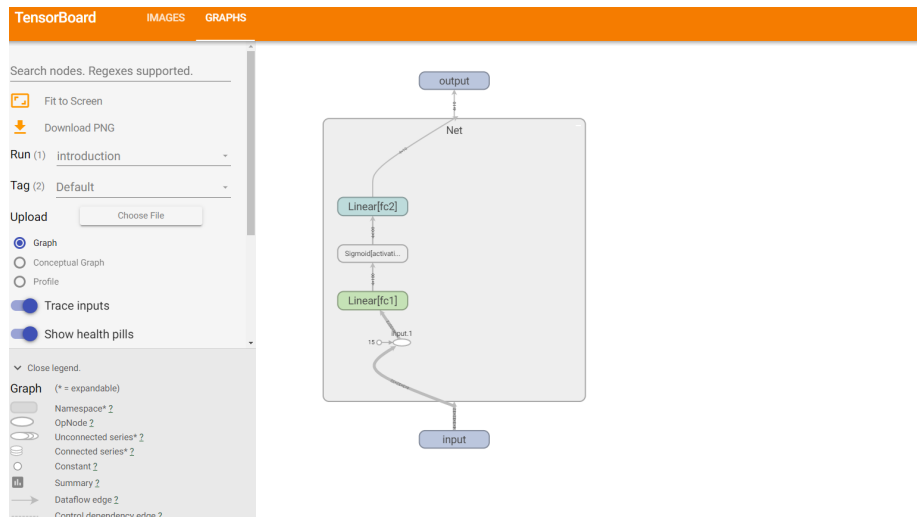
```
# Initialize your model
```

```
net = Net()
```

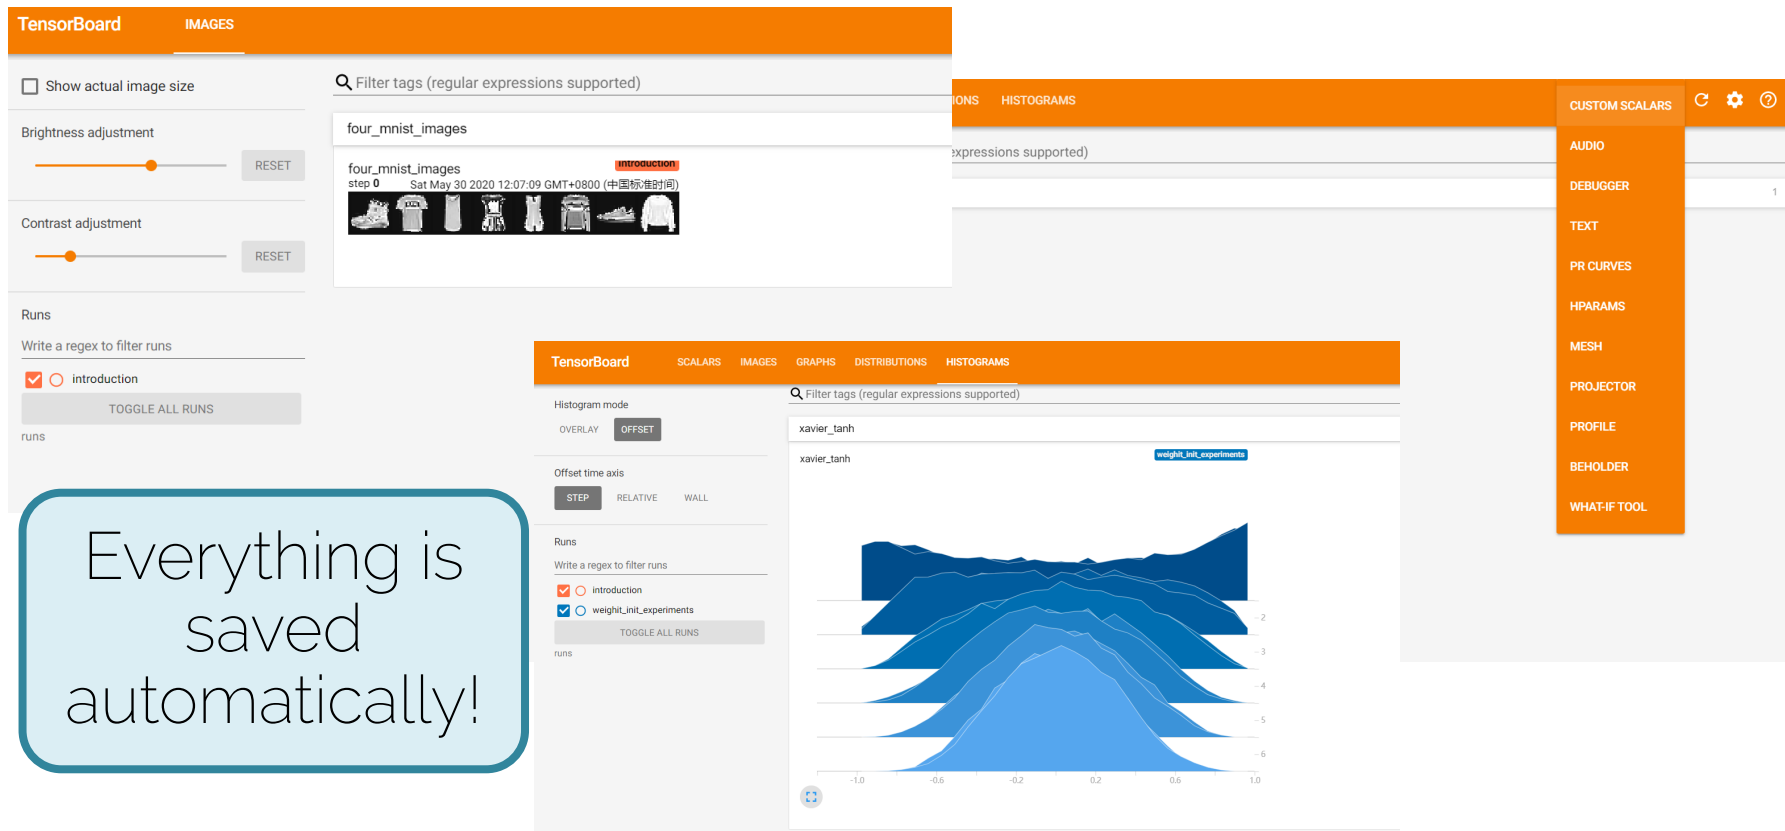
```
# Visualize its architecture in TensorBoard
```

```
writer.add_graph(net, images)
```

Graph creation
needs network
& one batch!



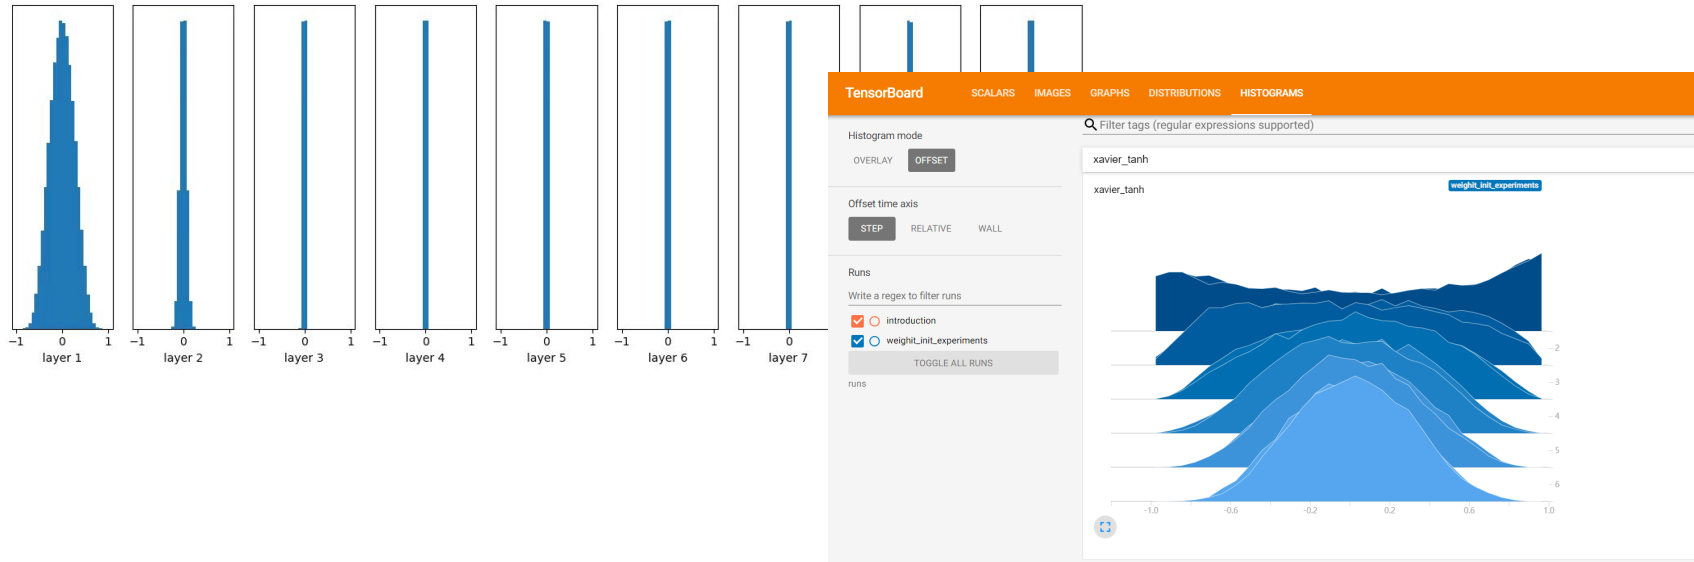
Document Everything!



Everything is
saved
automatically!

Weight Initialization

- Histogram visualization for layer outputs can show off effects of weight initialization as shown in the lecture



PyTorch Lightning

Idea Behind PyTorch Lightning

Classify our code into three categories

1. **Research** code (the exciting part!, changes with new tasks, models etc.)
2. **Engineering** code (the same for all projects and models)
3. **Non-essential** code (logging, organizing runs)

→ LightningModule

→ Trainer

→ Callbacks

Lightning Module

PyTorch

```
# model
class Net(nn.Module):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

# train loader
mnist_train = MNIST(os.getcwd(), train=True, download=True,
                    transform=transforms.ToTensor())
mnist_train = DataLoader(mnist_train, batch_size=64)

net = Net()

# optimizer + scheduler
optimizer = torch.optim.Adam(net.parameters(), lr=1e-3)
scheduler = SteplR(optimizer, step_size=1)

# train
for epoch in range(1, 100):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)

        loss.backward()
        optimizer.step()
        if batch_idx % args.log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'
                  .format(epoch, batch_idx * len(data), len(train_loader.dataset),
                          100. * batch_idx / len(train_loader), loss.item()))
```

Methods that need to be implemented

- `__init__`
- `forward`
- `training_step`
- `training_dataloader`
- `configure_optimizers`

Some other methods

- `validation_step`
- `validation_end`
- `prepare_data`

Trainer

PyTorch

PyTorch Lightning

```
# model
class Net(nn.Module):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

# train loader
mnist_train = MNIST(os.getcwd(), train=True, download=True,
                    transform=transforms.ToTensor())
mnist_train = DataLoader(mnist_train, batch_size=64)

net = Net()

# optimizer + scheduler
optimizer = torch.optim.Adam(net.parameters(), lr=1e-3)
scheduler = StepLR(optimizer, step_size=1)

# train
for epoch in range(1, 100):
    model.train()
    for batch_idx, (data, target) in enumerate(train_loader):
        data, target = data.to(device), target.to(device)
        optimizer.zero_grad()
        output = model(data)
        loss = F.nll_loss(output, target)

        loss.backward()
        optimizer.step()
        if batch_idx % args.log_interval == 0:
            print('Train Epoch: {} [{}/{} ({:.0f}%)]\tLoss: {:.6f}'.format(
                epoch, batch_idx * len(data), len(train_loader.dataset),
                100. * batch_idx / len(train_loader), loss.item()))
```

```
# model
class Net(LightningModule):
    def __init__(self):
        self.layer_1 = torch.nn.Linear(28 * 28, 128)
        self.layer_2 = torch.nn.Linear(128, 10)

    def forward(self, x):
        x = x.view(x.size(0), -1)
        x = self.layer_1(x)
        x = F.relu(x)
        x = self.layer_2(x)
        return x

    def train_dataloader(self):
        mnist_train = MNIST(os.getcwd(), train=True, download=True,
                            transform=transforms.ToTensor())
        return DataLoader(mnist_train, batch_size=64)

    def configure_optimizers(self):
        optimizer = torch.optim.Adam(self.parameters(), lr=1e-3)
        scheduler = StepLR(optimizer, step_size=1)
        return optimizer, scheduler

    def training_step(self, batch, batch_idx):
        data, target = batch
        output = self.forward(data)
        loss = F.nll_loss(output, target)
        return loss
```

```
if __name__ == '__main__':
    net = Net()
```

```
    trainer = Trainer()
    trainer.fit(net)
```

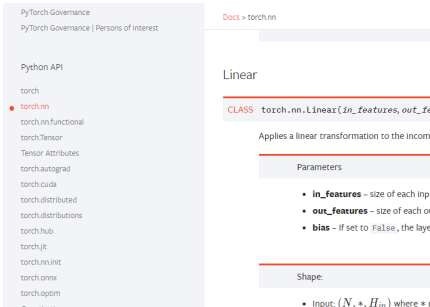
1. Initialize the model with hyperparameters for training (e.g. as a dictionary)
2. Trainer contains all code relevant for training our neural networks
3. Call the method `.fit()` for training the network

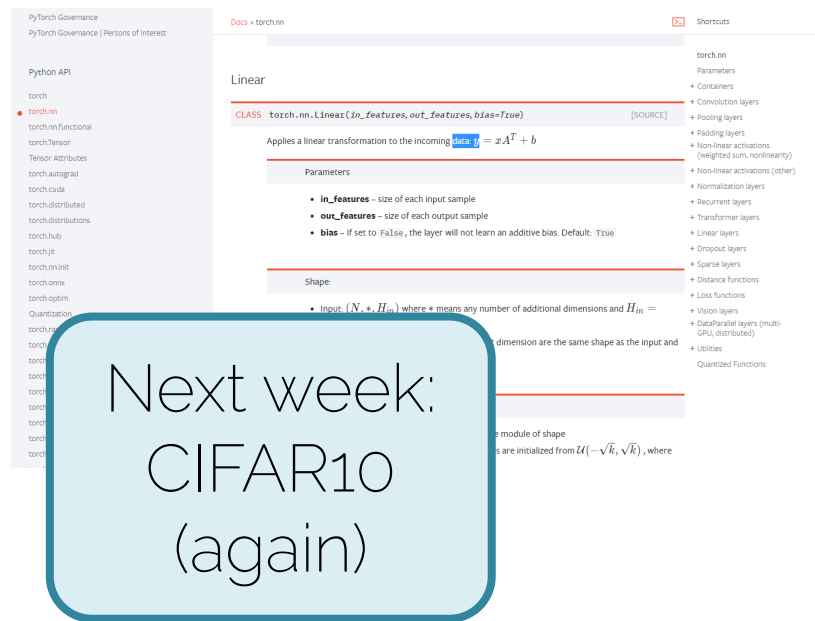
That's all you need to train you model 😊

Advantages

- Better overview of the relevant code
- Nice debugging features
- Many automated options, for example
 - Saving hyperparameters
 - Logging training/ validation results
 - Logging images with tensorboard
 - Saving models and reloading models from checkpoints

Your task for this week

- Go over all notebooks
 - There is no submission or code to implement
 - Experiment yourself!
 - House prices
 1. Pytorch with own solver
 2. Pytorch lightning
 - Try out visualizations
 - Check out documentations
- 
- The screenshot shows the PyTorch documentation for the `torch.nn.Linear` class. On the left, a sidebar lists various PyTorch components, with `torch.nn` expanded to show `torch.nn.Linear`. The main content area displays the class signature `torch.nn.Linear(in_features, out_features, bias=True)` and its description: "Applies a linear transformation to the incoming data: $y = Ax + b$ ". Below this, the "Parameters" section lists `in_features` (size of each input), `out_features` (size of each output), and `bias` (if set to `False`, the layer will not learn a bias vector). The "Shape:" section shows the input shape as `Input: (N, *, H_in)` where `*` is a variable length dimension.
- Next week:
CIFAR10



See you next week!