

# Exercise 3: Solution

# Dummy Dataset Forward Pass

```
31
32 class DummyNetwork(Network):
33     """
34     A Dummy network which takes in an input numpy array and computes its sigmoid
35     """
36
37     def __init__(self, model_name="dummy_model"): """
42
43     def forward(self, x):
44         """
45         :param x: The input to the network
46         :return: results of computation of sigmoid on the input
47         """
48         #####
49         # TODO
50         # Implement the sigmoid function.
51         #
52         #####
53         x = 1 / (1 + np.exp(-x))
54         #####
55         #                               END OF YOUR CODE
56         #####
57
58         return x
59
60     def __repr__(self):
61         return "A dummy class that would compute sigmoid function"
62
63     def save_model(self, data=None): """
69
```

# ImageFolderDataset make\_dataset()

```
40
41 @staticmethod
42 def make_dataset(directory, class_to_idx):
43     """
44     Create the image dataset by preparing a list of samples
45     :param directory: root directory of the dataset
46     :param class_to_idx: A dict that maps classes to labels
47     :returns: (images, labels) where:
48         - images is a list containing paths to all images in the dataset
49         - labels is a list containing one label per image
50     """
51     images, labels = [], []
52     #####
53     # TODO:
54     # Construct a list of all images in the dataset
55     # and a list of corresponding labels.
56     # Hints:
57     # - have a look at the "CIFAR-10: Image Dataset" notebook first
58     # - class_idx contains all classes and corresponding labels
59     # - images should only contain file paths, NOT the actual images
60     # - sort images by class and file name (ascending)
61     #####
62
63     for target_class in sorted(class_to_idx.keys()):
64         label = class_to_idx[target_class]
65         target_dir = os.path.join(directory, target_class)
66         for root, _, fnames in sorted(os.walk(target_dir)):
67             for fname in sorted(fnames):
68                 path = os.path.join(root, fname)
69                 images.append(path)
70                 labels.append(label)
71
72     #####
73     #                               END OF YOUR CODE
74     #####
75     assert len(images) == len(labels)
76     return images, labels
77
```

## Hints

- class\_to\_idx is a dictionary mapping the name of a label to an index.
- All images having the same label are found in a directory named just as the label.

## Code

- We iterate through the labels using class\_to\_idx.
- We obtain the path of the directory containing all images with a certain label (target\_dir) by concatenating the parent directory path (directory) with the name of the label (label).
- We get the names of all files in the target directory and add them to the images list.

# ImageFolderDataset \_\_len\_\_()

```
77
78 def __len__(self):
79     length = None
80     #####
81     # TODO:                                     #
82     # Return the length of the dataset (number of images) #
83     #####
84
85     length = len(self.images)
86
87     #####
88     #                                     END OF YOUR CODE                                     #
89     #####
90     return length
91
```

# ImageFolderDataset \_\_getitem\_\_()

## Hints

```
97 def __getitem__(self, index):
98     data_dict = None
99     #####
100     # TODO:
101     # create a dict of the data at the given index in your dataset
102     # The dict should be of the following format:
103     # {"image": <i-th image>, "label": <label of i-th image>}
104     # Hints:
105     # - use load_image_as_numpy() to load an image from a file path
106     # - Make sure to apply self.transform to the image
107     #####
108
109     label = self.labels[index]
110     path = self.images[index]
111     image = self.load_image_as_numpy(path)
112     if self.transform is not None:
113         image = self.transform(image)
114     data_dict = {
115         "image": image,
116         "label": label,
117     }
118
119     #####
120     #                               END OF YOUR CODE
121     #####
122     return data_dict
123
```

- self.images[index] contains the full name of the image we want to retrieve (we don't want to keep all images in memory at the same time - we only read them when it's required)
- Self.labels[index] contains the label of the image we want to retrieve
- We only apply the transformation if it's not None

# RescaleTransform \_\_call\_\_()

```
100
161 def __call__(self, images):
162     #####
163     # TODO:
164     # Rescale the given images:
165     #   - from (self._data_min, self._data_max)
166     #   - to (self.min, self.max)
167     #####
168
169     images = images - self._data_min # normalize to (0, data_max-data_min)
170     images /= (self._data_max - self._data_min) # normalize to (0, 1)
171     images *= (self.max - self.min) # norm to (0, target_max-target_min)
172     images += self.min # normalize to (target_min, target_max)
173
174     #####
175     #                               END OF YOUR CODE
176     #####
177     return images
178
```

# compute\_image\_mean\_and\_std()

```
124
125 def compute_image_mean_and_std(images):
126     """
127     Calculate the per-channel image mean and standard deviation of given images
128     :param images: numpy array of shape NxHxWxC
129     (for N images with C channels of spatial size HxW)
130     :returns: per-channels mean and std; numpy array of shape C
131     """
132     mean, std = None, None
133     #####
134     # TODO: #
135     # Calculate the per-channel mean and standard deviation of the images #
136     # Hint: You can use numpy to calculate mean and standard deviation #
137     #####
138
139     mean = np.mean(images, axis=(0, 1, 2))
140     std = np.std(images, axis=(0, 1, 2))
141
142     #####
143     #                               END OF YOUR CODE                               #
144     #####
145     return mean, std
146
```

# Dataloader: Iterator Helpers

```
26 def __iter__(self):
27     #####
28     # TODO:
29     # Define an iterable function that samples batches from the dataset.
30     # Each batch should be a dict containing numpy arrays of length
31     # batch size (except for the last batch if drop_last=True)
32     # Hints:
33     # - np.random.permutation(n) can be used to get a list of all
34     #   numbers from 0 to n-1 in a random order
35     # - To load data efficiently, you should try to load only those
36     #   samples from the dataset that are needed for the current batch.
37     #   An easy way to do this is to build a generator with the yield
38     #   keyword, see https://wiki.python.org/moin/Generators
39     # - Have a look at the "DataLoader" notebook first
40     #####
41     def combine_batch_dicts(batch):
42         """
43         Combines a given batch (list of dicts) to a dict of numpy arrays
44         :param batch: batch, list of dicts
45         e.g. [{k1: v1, k2: v2, ...}, {k1: v3, k2: v4, ...}, ...]
46         :returns: dict of numpy arrays
47         e.g. {k1: [v1, v3, ...], k2: [v2, v4, ...], ...}
48         """
49         batch_dict = {}
50         for data_dict in batch:
51             for key, value in data_dict.items():
52                 if key not in batch_dict:
53                     batch_dict[key] = []
54                 batch_dict[key].append(value)
55         return batch_dict
56
57     def batch_to_numpy(batch):
58         """Transform all values of the given batch dict to numpy arrays"""
59         numpy_batch = {}
60         for key, value in batch.items():
61             numpy_batch[key] = np.array(value)
62         return numpy_batch
```

## Hints

We create two helper functions: one for merging a batch of dictionaries as well as a convenient way to convert those dictionaries to numpy arrays which we will then feed to our networks later



# Dataloader: Iterator Implementation

```
63
64
65     if self.shuffle:
66         index_iterator = iter(np.random.permutation(len(self.dataset)))
67     else:
68         index_iterator = iter(range(len(self.dataset)))
69
70     batch = []
71     for index in index_iterator:
72         batch.append(self.dataset[index])
73         if len(batch) == self.batch_size:
74             yield batch_to_numpy(combine_batch_dicts(batch))
75             batch = []
76
77     if len(batch) > 0 and not self.drop_last:
78         yield batch_to_numpy(combine_batch_dicts(batch))
79
80     #####
81     #                               END OF YOUR CODE                               #
82     #####
```

## Hints

- Shuffling is implemented here using numpy's random permutation but there are multiple possible solutions
- We iterate over the dataset and use yield to properly invoke our iterator
- Finally we have to check for the last batch size in order to account for "drop\_last".

# Dataloader: Length

```
81
82     def __len__(self):
83         length = None
84         #####
85         # TODO:
86         # Return the length of the dataloader
87         # Hint: this is the number of batches you can sample from the dataset
88         #####
89         if self.drop_last:
90             length = len(self.dataset) // self.batch_size
91         else:
92             length = int(np.ceil(len(self.dataset) / self.batch_size))
93         #####
94         #                                     END OF YOUR CODE
95         #####
96         return length
97
```

Questions? Moodle 😊