

Exercise 5: Solution

Affine Layers

Affine Layer– Forward

```
def affine_forward(x, w, b):

    out = None
    #####
    # TODO: Implement the affine forward pass. Store the result in out. #
    # You will need to reshape the input into rows. #
    #####
    x_reshaped = np.reshape(x, (x.shape[0], -1))
    out = x_reshaped.dot(w) + b
    #####
    #                                END OF YOUR CODE                                #
    #####
    cache = (x, w, b)
    return out, cache
```

Remark: the input x , weights w and bias b are saved in cache, such that the backward pass can access them.

Affine Layer – Backward

```
def affine_backward(dout, cache):  
  
    x, w, b = cache  
    dx, dw, db = None, None, None  
    #####  
    # TODO: Implement the affine backward pass.                                #  
    # Hint: Don't forget to average the gradients dw and db                    #  
    #####  
    n = x.shape[0]  
  
    dw = (np.reshape(x, (x.shape[0], -1)).T).dot(dout) / n  
    dw = np.reshape(dw, w.shape)  
  
    db = np.mean(dout, axis=0, keepdims=False)  
  
    dx = dout.dot(w.T)  
    dx = np.reshape(dx, x.shape)  
    #####  
    #                                END OF YOUR CODE                                #  
    #####  
    return dx, dw, db
```

Remark:

Make sure the dw and dx have the same shape as w and x .

Here, we take the average of the gradient, because otherwise we have the sum over all gradients $\sum_{i=1}^N \nabla L_i(\theta)$ of the entire minibatch.

Hint: This averaging operation can also be done in the backward pass of the loss function. If so, we don't average n in the layer.

Sigmoid – Forward

```
def sigmoid_forward(x):  
    """  
    Computes the forward pass for a layer of sigmoids.  
  
    :param x: Inputs, of any shape  
  
    :return out: Output, of the same shape as x  
    :return cache: out  
    """  
    out = None  
    #####  
    # TODO: Implement the Sigmoid forward pass. #  
    #####  
    out = 1 / (1 + np.exp(-x))  
    #####  
    #                                     END OF YOUR CODE #  
    #####  
    cache = out  
    return out, cache
```

Remark:

The output of sigmoid function is stored in the cache for the computation in backward pass.

Sigmoid – Backward

```
def sigmoid_backward(dout, cache):
    """
    Computes the backward pass for a layer of sigmoids.

    :param dout: Upstream derivatives, of any shape
    :param cache: y, output of the forward pass, of same shape as dout

    :return dx: Gradient with respect to x
    """
    dx = None
    y = cache
    #####
    # TODO: Implement the Sigmoid backward pass. #
    #####
    dx = dout * y * (1 - y)
    #####
    #                               END OF YOUR CODE                               #
    #####
    return dx
```

Remark:

The derivative of sigmoid function is
is $\text{sigmoid} * (1 - \text{sigmoid})$

Two Layer Network

Forward Pass

```
#####  
# TODO  
# Implement the forward pass using the layers you implemented.  
# It consists of 3 steps:  
# 1. Forward the first affine layer  
# 2. Forward the sigmoid layer  
# 3. Forward the second affine layer  
# (Dont't forget the caches)  
#####  
  
# Forward first layer  
h, cache_affine1 = affine_forward(X, W1, b1)  
  
# Activation function  
h_, cache_sigmoid = sigmoid_forward(h)  
  
# Forward second layer  
y, cache_affine2 = affine_forward(h_, W2, b2)  
  
#####  
#                               END OF YOUR CODE  
#####
```

Remark:

The weights and biases are initialized in `__init__` of the class.

The first affine layer takes the input X , weights $W1$, bias $b1$, and returns the output h for the input of next layer and cache *cache_affine1* for later gradient computation.

Same procedure applies to other layers.

Backward Pass

```
#####  
# TODO                                                                    #  
# Implement the backward pass using the layers you implemented.          #  
# Like the forward pass, it consists of 3 steps:                        #  
#   1. Backward the second affine layer                                  #  
#   2. Backward the sigmoid layer                                         #  
#   3. Backward the first affine layer                                    #  
# You should now have the gradients wrt all model parameters            #  
#####  
  
# Backward second layer  
dh_, dw2, db2 = affine_backward(dy, cache_affine2)  
  
# Backward activation function  
dh = sigmoid_backward(dh_, cache_sigmoid)  
  
# Backward first layer  
dx, dw1, db1 = affine_backward(dh, cache_affine1)  
  
#####  
#                               END OF YOUR CODE                          #  
#####
```

Review of Exercise 5

https://docs.google.com/forms/d/e/1FAIpQLSdfENBvUMsnyXkO7Bf_VINaUgTrEv5EBWYMGhq8Q77tCrrjqg/viewform?usp=sf_link

Questions? Moodle

