

Exercise 8: Solution

Disclaimer

- In this solution, we just present a minimal model that passes the threshold
- Feel free to play around with different architectures and come up with your own model!
- But also note that you don't always need hundreds of millions parameters to solve a simple problem like this one 😊

Define your Hyperparameters

```
from exercise_code.MyPytorchModel import MyPytorchModel

hparams = {}
#####
# TODO: Define your hyper parameters here! #
#####
hparams = {
    "batch_size": 64,
    "learning_rate": 3e-4
}
#####
#                               END OF YOUR CODE                               #
#####
model = MyPytorchModel(hparams)
model.prepare_data()
```

Define your Trainer

```
trainer = None

#####
# TODO: Define your trainer!                                     #
#####

trainer = pl.Trainer(
    max_epochs = 10,
    gpus = 1 if torch.cuda.is_available() else None
)

#####
#                               END OF YOUR CODE                #
#####

trainer.fit(model)
```

Initialize your Model

```
def __init__(self, hparams, input_size=3 * 32 * 32, num_classes=10):
    super().__init__()

    # set hyperparams
    self.hparams = hparams
    self.model = None

    #####
    # TODO: Initialize your model!                                #
    #####

    self.model = nn.Sequential(
        nn.Linear(input_size, 500),
        nn.BatchNorm1d(500),
        nn.ReLU(),
        nn.Dropout(p=0.5),
        nn.Linear(500, 100),
        nn.BatchNorm1d(100),
        nn.ReLU(),
        nn.Dropout(p=0.5),
        nn.Linear(100, 10)
    )

    #####
    #                                END OF YOUR CODE                                #
    #####
```

Just an example.
Be creative here and come up
with your own architecture 😊

Prepare Data

```
#####
# TODO: Define your transforms (convert to tensors, normalize).      #
# If you want, you can also perform data augmentation!              #
#####
mean = [0.485, 0.456, 0.406]
std = [0.229, 0.224, 0.225]
my_transform = transforms.Compose([
    transforms.RandomApply((transforms.RandomHorizontalFlip(p=0.8),
                           transforms.RandomResizedCrop((32,32))), p=0.1),
    transforms.ToTensor(),
    transforms.Normalize(mean, std)
])
#####
#                               END OF YOUR CODE                      #
#####
cifar_complete = torchvision.datasets.ImageFolder(root=CIFAR_ROOT, transform=my_transform)
```

Define Optimizers

```
def configure_optimizers(self):

    optim = None

    #####
    # TODO: Define your optimizer.                                     #
    #####

    optim = torch.optim.Adam(self.model.parameters(), self.hparams["learning_rate"])

    #####
    #                               END OF YOUR CODE                               #
    #####

    return optim
```

Questions? Moodle

